

# Artificial Intelligence

AI-2002

Lecture 6

Mahzaib Younas

Lecturer Department of Computer Science

FAST NUCES CFD

# Beyond Classical Search

# Beyond Classical Search


- We have addressed a single category of problems:  
**observable, deterministic, known environments**
  - where the solution is a sequence of actions
- The search algorithms that we have seen so far are designed to explore search spaces **systematically**.
  - When a **goal is found**, the **path** to that goal also constitutes a **solution** to the problem.

# Beyond Classical Search

- In many problems, however, the **path to the goal is irrelevant**.

## For example:

- 8-queens problem
  - *what matter, is the final configuration of queens, not the order.*
- The factory-floor layout problem
- The vehicle routing problem



# Need Of Local Search

- The first issue is that
  - The algorithm tries to explore the entire search space systematically.
  - The algorithm visits the states in a certain order and it may visit a lot of the states before finding a goal.
  - There are some obvious problems with this behaviour. If the search space is big, systematic exploration will take a long time.
  - If the search space is infinite, we cannot hope to visit all of states
  - the search algorithm remembers and returns a path from the initial state to the goal state

# Properties Of Local Search (Cont...)

- First, local search does not attempt to explore the search space systematically.
- Second, local search only remembers the current state and does not keep track of a path to the goal node

# Local Search Algorithms

# Local Search Algorithm

local search algorithms **give up on exploring the search space systematically.**

instead of attempting to visit all of the states, local search uses **strategies to find reasonably good states quickly on average**

good enough in practice especially when we are **solving a challenging problem under time constraint.**



# Local Search Algorithm

A local search problem consists of:

- A state
  - a complete assignment to all of the variables.
- A neighbour relation
  - which states do I explore next?
- A cost function
  - how good is each state?

# Local Search Algorithms

- Operate using a **single current node** and generally move only to neighbors of that node.
  - *No concern with paths* followed by the search
  - They are *NOT systematic*
- Local search algorithms have **two key advantages**:
  - they use very *little memory*
  - they can often find reasonable *solutions in large* or *infinite (continuous)* state spaces

# Local Search Algorithms

- The local search algorithms are useful for solving pure **optimization problems**,

## Optimization problems:


- the aim is **to find the best state** according to an *objective function*.
- The *standard search methods do not fit well* with optimization problems.

# Local Search Algorithms

## State-space landscape

A landscape has both

- “**location**”
  - defined by the state
- “**elevation**”
  - defined by the value of the **heuristic cost function** or **objective function**.




# Local Search Algorithms

## State-space landscape

- If elevation corresponds to cost, then the aim is to find the lowest valley—a **global minimum**;
- If elevation corresponds to an **objective function**, then the aim is to find the highest peak—a **global maximum**.

# Local Search Algorithms

## State-space landscape



A **one-dimensional state-space landscape** in which *elevation corresponds to the objective function*. The aim is to find the **global maximum**.

# Local Search Algorithms

- A **complete local search** algorithm always finds a goal if one exists;
- An **optimal** algorithm always finds a global minimum/maximum.
- Local search algorithms typically use a **complete-state formulation**,
  - In 8-queen problem, where *each state has 8 queens* on the board, one per column.


# Terminology

<u>Terminology</u>	<u>Purpose</u>
<b>Shoulder</b>	It is a region having an edge upwards and it is also considered as one of the problems in hill climbing algorithms.
<b>Global Maximum</b>	a state that maximizes the objective function over the entire landscape.
<b>Local Maximum</b>	it is the state which is slightly better than the neighbor states but it is always lower than the highest state.
<b>Flat Maximum</b>	if the neighbor states all having same value, they can be represented by a flat space (as seen from the diagram) which are known as flat local maximums.
<b>Current state</b>	It is the state which contains the presence of an active agent.




# Example: $n$ -queens Problems


- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts
- Heuristic  $h$ : number of ‘attacks’



$h = 5$



$h = 2$






$h = 0$

# Example: $n$ -queens Problems

- Put  $n$  queens on an  $n \times n$  board with no two queens on the same row, column, or diagonal
- Move a queen to reduce number of conflicts
- Heuristic  $h$ : number of ‘attacks’

If queen is placed here, there is only one attack then,



# Hill-climbing Search

# Hill-climbing Search

- It is simply a **loop** that continuously moves in the direction of increasing value—that is, uphill.
- It terminates when it reaches a “**peak**” where no neighbor has a higher value.
  - *does not maintain a search tree*
  - need only  
*record the state* and  
the value of the *objective function*.
- Hill climbing only looks towards the *immediate neighbors* of the current state.

# Hill-climbing Search

- The hill climbing often gets stuck for the following reasons:

## **Local maxima:**

- A local maximum is a peak that is higher than each of its neighboring states but lower than the global maximum.


# Hill-climbing Search

## Plateaux:

- An area of the state space where the **evaluation function is flat**.
- It can be a **flat local maximum**, from which no uphill exit exists,

## Shoulder:


- from which progress is possible




# Hill-climbing Search

## Ridges

- A sequence of local maxima
- *Its difficult for greedy algorithms to navigate*



# Hill-climbing Search





# Hill-climbing Search

**function** HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

*current*  $\leftarrow$  MAKE-NODE(*problem*.INITIAL-STATE)

**loop do**

*neighbor*  $\leftarrow$  a highest-valued successor of *current*

**if** *neighbor*.VALUE  $\leq$  *current*.VALUE **then return** *current*.STATE

*current*  $\leftarrow$  *neighbor*

- At *each step the current node is replaced by the best neighbor*; the neighbor with the highest VALUE,
- If a **heuristic cost estimate *h*** is used, we would find the neighbor with the *lowest h*.

# Hill Climb Search

- It is simply a **loop** that continuously moves in the direction of increasing value—that is, uphill.
- It terminates when it reaches a “**peak**” where no
- neighbor has a higher value.
  - **does not maintain a search tree**
  - need only
    - record the state and
    - the value of the **objective function**.
- Hill climbing only looks towards the immediate neighbors of the current state.

# Hill Climb Search (Example)

- With randomly generated 8-queens starting states, the steepest-ascent hill climbing:
  - 14% of the time it solves the problem
  - 86% of the time it get stuck at a local minimum

# Types/Variants of Hill Climb

- **Stochastic hill-climbing:**
  - Chooses randomly among potential successors
  - Sometimes better than steepest ascent
- **First-choice hill-climbing:**
  - Generates successors randomly and picks first
  - Good for many successors
- **Random restart hill-climbing:**
  - Restarts from randomly generated initial state when failed
  - Roughly 7 iterations with 8-queens problem

# Hill Climb Search (Cont...)

- The success of hill-climbing search depends very much
  - on the shape of the state-space landscape
- If there are few local maxima and plateaux,
  - random-restart hill climbing will find a good solution very quickly.

# Local Beam Search

# Local Beam Search

## Idea:

- □ Keep track of  $k$  states rather than just one
- □ Start with  $k$  randomly generated states
- □ At each iteration, all the successors of all  $k$  states are generated
- □ If any one is a goal state, stop;
- □ else select the  $k$  best successors from the complete list and repeat.

# Local Beam Search

```
function BEAM-SEARCH(problem, k) returns a solution state
  start with k randomly generated states
  loop
    generate all successors of all k states
    if any of them is a solution then return it
    else select the k best successors
```

A local beam search with  $k$  states might seem to be nothing more than running  **$k$  random restarts in parallel** instead of in sequence.



# Local Beam Search

## Local beam search with $k = 1$

- We would randomly generate **1 start state**
- At each step we would generate all the successors, and retain the **1 best state**
- Equivalent to **HILL-CLIMBING**

# Local Beam Search

## Local beam search with $k = \infty$

**1 initial state** and no limit of the number of states retained

We start at initial state and **generate ALL successor** states (no limit how many)

If one of those is a goal, we stop

Otherwise, we generate all successors of those states (2 steps from the initial state), and continue

Equivalent to **BREADTH-FIRST SEARCH** except that

*each layer is generated all at once.*


# Hill Climbing Vs. Beam Search

- Hill climbing just explores all nodes in one branch until goal found or not being able to explore more nodes.
- Beam search explores more than one path together. A factor **k** is used to determine the number of branches explored at a time.
- If k=2, then two branches are explored at a time. For k=4, four branches are explored simultaneously.
- The branches selected are the **best branches** based on the used **heuristic evaluation function**.

# Beam Search, $k=2$

## Goal – Node K

Current	Children
a	---



# Beam Search


## Goal – Node K

Current

a

Children


---



# Beam Search


## Goal – Node K

Current	Children
<b>a</b>	---
a	



# Beam Search

## Goal – Node K



Current

a

a

Children


---

*f7, j8, b10*

# Beam Search

## Goal – Node K

Best k Successors	Current	Children
	a	---
	a	<i>f7, j8, b10</i>






# Beam Search


## Goal – Node K

Best k Successors	Current	Children
	a	---
	a	<i>f</i> 7, <i>j</i> 8, <i>b</i> 10



# Beam Search


## Goal – Node K



Current		Children
a		---
a		<b>f7, j8, b10</b>
f	j	

# Beam Search


## Goal – Node K



Current	Children
a	---
a	<b>f7, j8, b10</b>
<u>f</u> <u>j</u>	

# Beam Search

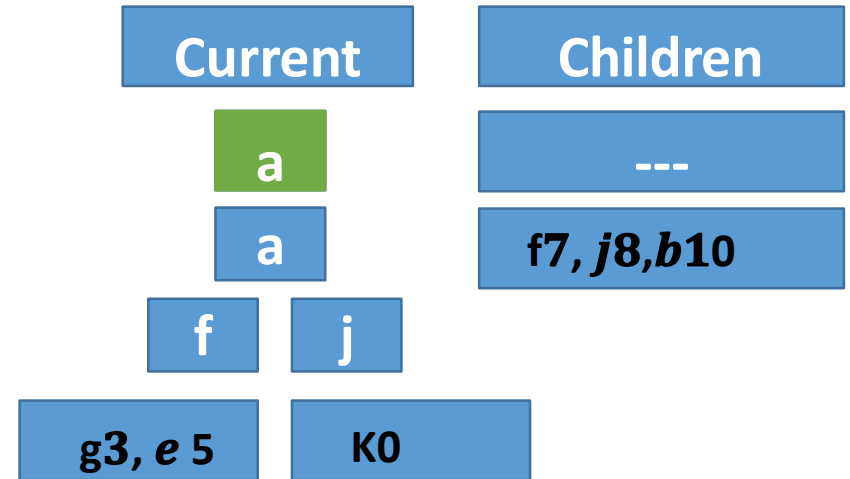

## Goal – Node K



Current	Children
a	---
a	<b>f7, j8, b10</b>
f	
j	

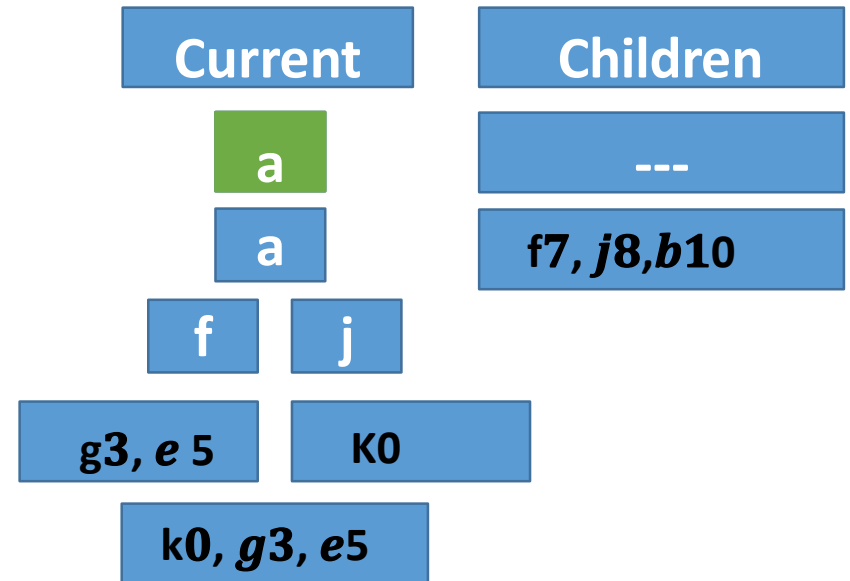

# Beam Search

## Goal – Node K




# Beam Search

## Goal – Node K

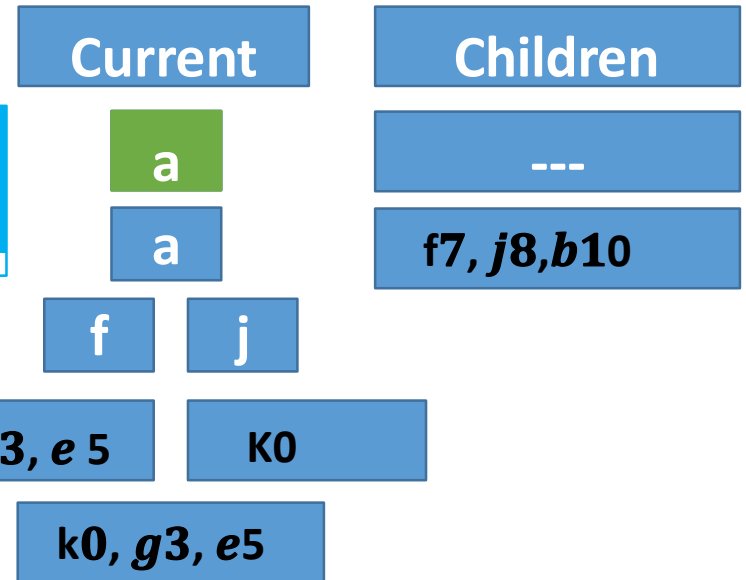


# Beam Search

## Goal – Node K

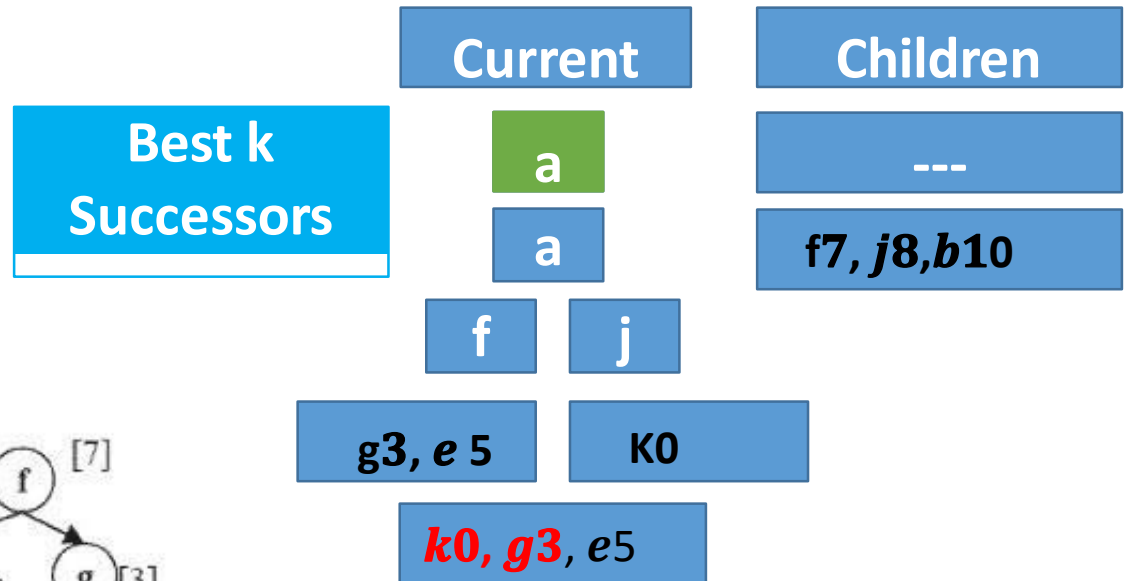



**Best k  
Successors**



# Beam Search

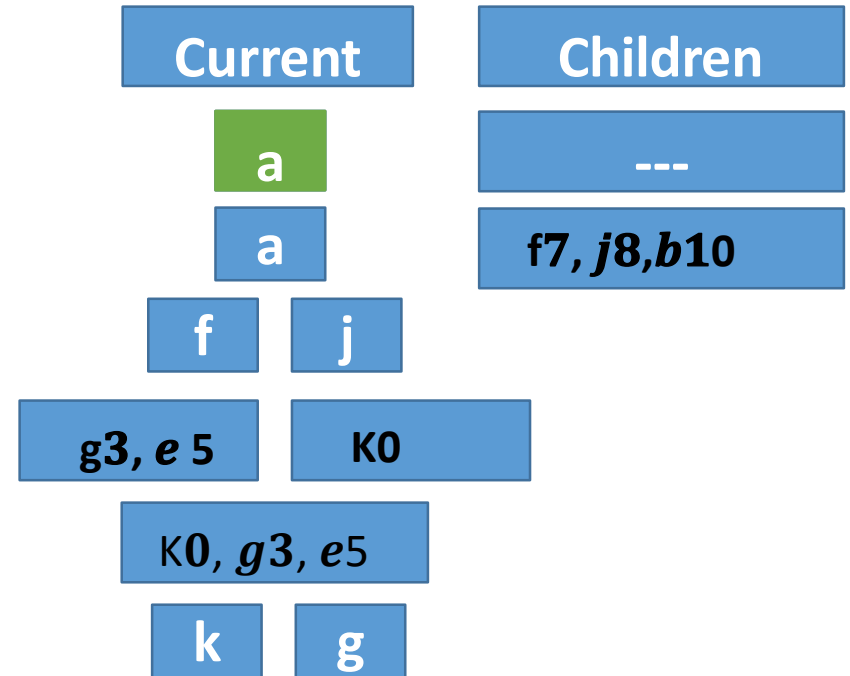

## Goal – Node K





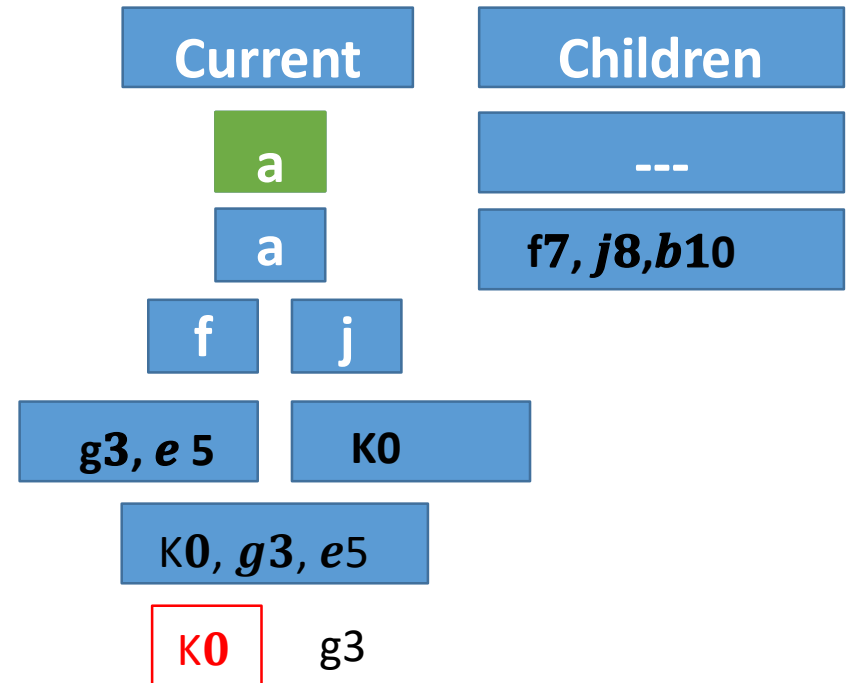

# Beam Search

## Goal – Node K



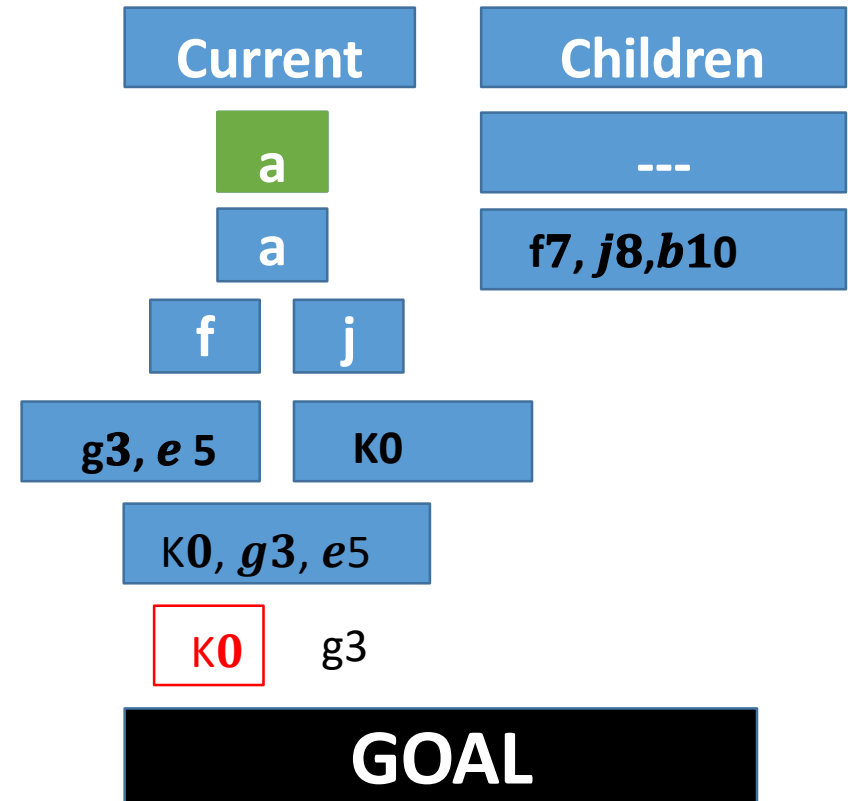

# Beam Search

## Goal – Node K



# Beam Search

## Goal – Node K



# **Simulated Annealing**

# Simulated Annealing

- Idea:
  - escape local maxima by allowing some “bad” moves but gradually decrease the size and frequency of the bad moves,
- In thermodynamics, the probability to go from a state with energy  $E_1$  to a state of energy  $E_2$  is given by:

$$p = e^{\frac{(E_2 - E_1)}{kT}} = e^{\frac{-(E_1 - E_2)}{kT}}$$

# Simulated Annealing (cont...)

- $e$  is Euler's number
- $T$  is a "temperature" controlling the probability of downward steps
- $k$  is Boltzmann's constant
  - (relating energy and temperature; with appropriate choice of units, it will be equal to 1).

# Simulated Annealing

$$p = e^{\frac{-(E_1 - E_2)}{kT}}$$

Where,

- $e$  is Euler's number
- $T$  is a "temperature" controlling the probability of downward steps
- $k$  is Boltzmann's constant
  - (relating energy and temperature; with appropriate choice of units, it will be equal to 1).

# Simulated Annealing

$$p = e^{\frac{(E_2 - E_1)}{kT}} = e^{\frac{-(E_1 - E_2)}{kT}}$$

The idea is that probability decreases exponentially with  $E_2 - E_1$  increasing,

The probability gets lower as temperature decreases

If the *schedule* lowers  $T$  slowly enough, the algorithm will find a global optimum with probability approaching 1.



# Simulated Annealing

function **SIMULATED-ANNEALING**(*problem*, *schedule*) returns a solution state

inputs: *problem*, a problem

*schedule*, a mapping from time to “temperature”

local variables: *current*, a node

*next*, a node

*T*, a “temperature” controlling prob. of downward steps

*current* ← MAKE-NODE(INITIAL-STATE[*problem*])

for *t* ← 1 to  $\infty$  do

*T* ← *schedule*[*t*]

if *T* = 0 then return *current*

*next* ← a randomly selected successor of *current*

$\Delta E$  ← VALUE[*next*] – VALUE[*current*]

if  $\Delta E > 0$  then *current* ← *next*

else *current* ← *next* only with probability  $e^{\Delta E/T}$

Similar to hill climbing,  
but a random move instead  
of best move

case of improvement, make  
the move

Otherwise, choose the move with probability that  
decreases exponentially with the “badness” of the move.

# Simulated Annealing...Example

Consider there are 3 moves available, with changes in the objective function of

$$\Delta E_1 = -0.1, \Delta E_2 = 0.5, \Delta E_3 = -3$$

Suppose  $T = 1$

Pick a move randomly:

- if  $\Delta E_2$  is picked, move there.
- if  $\Delta E_1$  or  $\Delta E_3$  are picked, probability of move =  $e^{\frac{\Delta E}{T}}$ 
  - move 1:  $\text{prob1} = e^{-0.1} = 0.9$ ,  
i.e., 90% of the time we will accept this move
  - move 3:  $\text{prob3} = e^{-3} = 0.0497$   
i.e., 5% of the time we will accept this move

# Simulated Annealing

$T$  = “temperature” parameter

**If  $T$  is high**  $\Rightarrow$  the probability of “locally bad” move is higher

**If  $T$  is low**  $\Rightarrow$  the probability of “locally bad” move is lower

typically,  $T$  is decreased as the algorithm runs longer

- i.e., there is a “temperature schedule”

# Simulated Annealing

Convergence:

With exponential schedule, will provably converge to global optimum

- If  $T$  decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1.

Few more precise convergence rate

Recent work on **rapidly mixing Markov chains**.

Surprisingly, deep foundations.

# Simulated Annealing

method proposed in 1983 by IBM researchers for solving VLSI layout problems.

- theoretically will always find the global optimum (the best solution)

Useful for some problems, but can be very **slow**

- slowness comes about because  $T$  must be decreased very gradually to retain optimality

In practice how to decide the rate at which to decrease

$T$ ? (this is a practical problem with this method)