

# Artificial Intelligence

## AI-2002

### Lecture 8

Mahzaib Younas  
Lecturer Department of Computer Science  
FAST NUCES CFD

# Adversarial Search

- **Competitive** environments, in which the **agents' goals are in conflict**, giving rise to **adversarial search** problems—often known as **games**

## Why do AI researchers study game playing?

- It's a **good reasoning problem**, formal and nontrivial.
- Offer an opportunity to study problems involving **{hostile, adversarial, competing} agents**.
- **Direct comparison** with humans and other computer programs is easy.

# Adversarial Search

Mainly games of strategy with the following characteristics:

- Sequence of **moves** to play
- Rules that specify **possible moves**
- Rules that specify a **payment** for each move
- Objective is to **maximize** your payment

# Games

- **Competitive: Commonly Zero Sum**
  - One player wins and the other loses
  - A zero-sum game is defined as one where the **total payoff to all players is the same for every instance** of the game. Chess is zero-sum because every game has payoff of either  $0 + 1$ ,  $1 + 0$  or  $\frac{1}{2}$ ,  $\frac{1}{2}$ .
- **Perfect Information:**
  - Players know the results of all previous moves
  - There is one best way to win the game for all players
- **Imperfect Information**
  - Players do not know all of the previous moves

# Games

- **Initial State ( $s_0$ ):** The initial state, which specifies how the game is *set up at the start*.
- **Players:** defines which player has the move in a state.
- **Actions:** The set of legal moves.
- **Result ( $s, a$ ):** The transition model, which defines the result of a move. The state after action  $a$  is the state  $s$ .
- **Terminal Test:** A terminal test, which is true when the game is over and false otherwise.


# Games

- **Terminal State:** States where the game has ended are called terminal states.
- **Utility:** A utility function (also called an **objective function** or **payoff function**), defines the final numeric value for a game that ends in terminal state ***s*** for a player ***p***.
  - In chess, the outcome is a win, loss, or draw, with values +1, 0, or ½.
  - Some games have a wider variety of possible outcomes; the payoffs in **backgammon range from 0 to +192**.

# Game Tree

- The **initial state**, **actions**, and **results** define the game tree for the game.
- A **game tree** where the **nodes are the game states** and **the edges are moves**.
- The game tree is best thought of as a theoretical construct that we cannot realize in the physical world.
  - For **tic-tac-toe** the game tree is relatively small—fewer than  $9! = 362,880$  terminal nodes.
  - For **chess** there are over  $10^{40}$  nodes,

# Game Tree: tic-tac-toe





# Minimax

- Minimax is a method used **to evaluate game trees**.
- Given a game tree, the optimal strategy can be determined from the **minimax value** of each node.
- *A static evaluator is applied to leaf nodes*, and **the values are passed back up the tree to determine the best score** the computer can obtain against a rational opponent.

# Minimax


## MAX

- Wants to **maximize** the result of the utility function
- **Winning strategy** if, on MIN's turn, a win is obtainable for MAX for all moves that MIN can make

## MIN

- Wants to **minimize** the result of the utility function
- **Winning strategy** if, on MAX's turn, a win is obtainable for MIN for all moves that MAX can make


# Minimax



MINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$


# Minimax



MINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

# Minimax



MINIMAX( $s$ ) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

# Minimax

**function** MINIMAX-DECISION(*state*) *returns an action*  
    **return**  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$

---

**function** MAX-VALUE(*state*) *returns a utility value*  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow -\infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
    **return** *v*

---

**function** MIN-VALUE(*state*) *returns a utility value*  
    **if** TERMINAL-TEST(*state*) **then return** UTILITY(*state*)  
     $v \leftarrow \infty$   
    **for each** *a* **in** ACTIONS(*state*) **do**  
         $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
    **return** *v*

# Properties of Minimax

## Complete?

- Yes (if tree is finite).

## Optimal?

- Yes

## Time complexity?

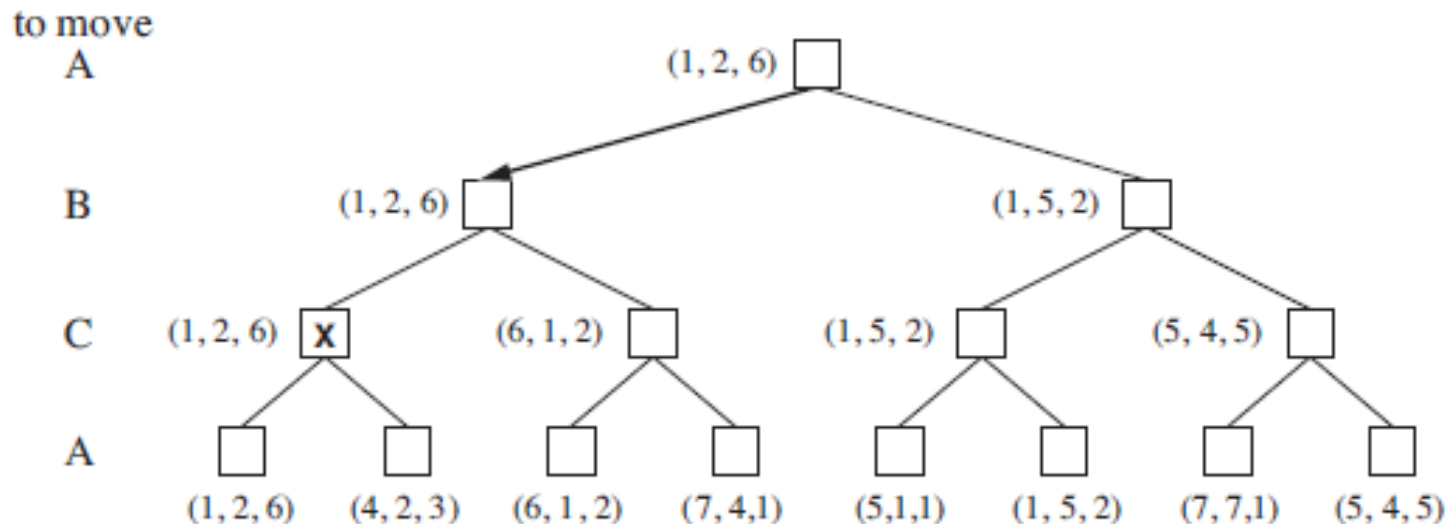
- $O(b^m)$ ,  $m$  is the maximum depth of the tree and  $b$  is the legal moves.

## Space complexity?

- $O(bm)$ 
  - (depth-first search, generate all actions at once)
- $O(m)$ 
  - (backtracking search, generate actions one at a time)

# Multiplayer Games

- Each node must hold **a vector of values**
- For example, for three players  $A, B, C$  ( $V_A, V_B, V_C$ )
- The backed up vector at node  $n$  will always be the one that **maximizes the payoff** of the player choosing at  $n$






# Searching Game Trees

- **Exhaustively searching** a game tree is not usually a good idea.
- Even for a simple game like
  - **tic-tac-toe** there are over **350,000 nodes** in the complete game tree.
- **An additional problem** is that the computer only gets to choose every other path through the tree and the opponent chooses the others.

# Pruning

MAX

MIN




$$\begin{aligned} \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\ &= \max(3, \min(2, x, y), 2) \end{aligned}$$

# Pruning

MAX

MIN



$$\begin{aligned}
 \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\
 &= \max(3, \min(2, x, y), 2) \\
 &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\
 &= 3.
 \end{aligned}$$

**Do we need  $z$ ?**

# Pruning

- We can use a **branch-and-bound technique** to reduce the number of states that must be examined to determine the value of a tree.

## Branch-and-bound Technique:

- We keep track of a **lower bound on the value of a maximizing node**, and don't bother evaluating any trees that cannot improve this bound.
- Keep track of an **upper bound on the value of a minimizing node**. Don't bother with any sub-trees that cannot improve this bound.

# Minimax with Alpha-Beta Cutoffs

## Alpha Cutoffs:

- Alpha is the *lower bound on maximizing nodes*.


## Beta Cutoffs:

- Beta is the *upper bound on minimizing nodes*.
- Both alpha and beta get passed down the tree during the Minimax search.


# Minimax with Alpha-Beta Cutoffs

- At minimizing nodes, we stop evaluating children if we get a child whose value is **less than the current lower bound (*alpha*)**.
- At maximizing nodes, we stop evaluating children as soon as we get a child whose value is **greater than the current upper bound (*beta*)**.
- Some branches will never be played by rational players since they include sub-optimal decisions (for either player)

# Alpha-Beta Pruning: Example




# Alpha-Beta Pruning: Example






# Alpha-Beta Pruning: Example




# Alpha-Beta Pruning: Example



# Alpha-Beta Pruning: Example

At MIN we will change  
the value of BETA


At MAX we will change  
the value of ALPHA



# Alpha-Beta Pruning: Example

At MIN we will change  
the value of BETA


At MAX we will change  
the value of ALPHA



# Alpha-Beta Pruning: Example

At MIN we will change  
the value of BETA

At MAX we will change  
the value of ALPHA



# Alpha-Beta Pruning: Effectiveness

- The effectiveness **depends on the order** in which children are visited.
- **In the best case**, the effective branching factor will be reduced from  **$b$**  to  **$\sqrt{b}$** .
- **In an average case** (random values of leaves) the branching factor is reduced to  $\frac{b}{\log b}$ .

# Reading Material

- **Artificial Intelligence, A Modern Approach**  
**Stuart J. Russell and Peter Norvig**
  - **Chapter 5.**